

AI Infrastructure & SLA Report

LLM Token Consumption, Latencies & Errors (12/8/2026 - 26/8/2026)

BE-08 AGY

AI Core Metrics

TOTAL API CALLS

135

TOTAL COST

\$1.69

AVG LATENCY

1201 ms

UPTIME SLA

79.26%

Model Utilization Details

Model	Calls	Prompt Tokens	Comp Tokens	Avg Latency	Cost
llama-3-70b	38	70,176	32,035	1212 ms	\$0.1202
gpt-4o	35	54,383	27,316	1181 ms	\$0.6816
claude-3-5-sonnet	33	49,314	32,287	1273 ms	\$0.4462
gemini-1-5-pro	29	53,420	24,399	1129 ms	\$0.4432

API HTTP Response Code Distribution

Status Code	Status Type	Total Responses	Percentage
200	SUCCESS (200 OK)	107	79.26%
429	RATE LIMIT (429)	16	11.85%
500	SERVER ERROR (500)	12	8.89%

