

AI Infrastructure & SLA Report

LLM Token Consumption, Latencies & Errors (19/8/2026 - 26/8/2026)

BE-08 AGY

AI Core Metrics

TOTAL API CALLS

60

TOTAL COST

\$0.69

AVG LATENCY

1209 ms

UPTIME SLA

80%

Model Utilization Details

Model	Calls	Prompt Tokens	Comp Tokens	Avg Latency	Cost
llama-3-70b	20	38,041	18,352	1232 ms	\$0.0672
gpt-4o	16	25,474	11,088	1167 ms	\$0.2937
claude-3-5-sonnet	14	18,059	13,533	1129 ms	\$0.1805
gemini-1-5-pro	10	18,220	7,997	1342 ms	\$0.1477

API HTTP Response Code Distribution

Status Code	Status Type	Total Responses	Percentage
200	SUCCESS (200 OK)	48	80.00%
500	SERVER ERROR (500)	6	10.00%
429	RATE LIMIT (429)	6	10.00%

